

Addressing Classification challenges in Fine-Grained Political Evasion Detection via LLM-Based Data Manipulation

Marta Leš, Ljubomir Jerončić, Luka Špiljak

University of Zagreb, Faculty of Electrical Engineering and Computing
Unska 3, 10000 Zagreb, Croatia
{marta.les, ljubomir.jeroncic, luka.spiljak}@fer.hr,

Abstract

This paper presents our approach to the SemEval-2026 CLARITY shared task (Task 6) on political question evasion detection, addressing both Task 1 (three-class clarity classification) and Task 2 (nine-class fine-grained evasion classification) using question-answer pairs from U.S. presidential interviews. We fine-tune RoBERTa-base as our encoder baseline and investigate three research directions: (1) an error analysis of confusion patterns between semantically similar evasion labels, revealing clusters that confirm a "meaning overlap" between classes effecting model failure, (2) LLM-based data augmentation targeting minority evasion classes using a GPT model pipeline including generation and quality filtering of generated data, as well as replacing the given GPT-3.5 summary features with new GPT-5.4-nano generated ones, last (3) is a three-stage confidence-conditional prompt-based routing pipeline that hierarchically decomposes the nine-class problem into clarity-level routing followed by subset-specific evasion classification. Our best system achieves a macro F1 of 0.73 on Task 1 and 0.53 on Task 2 using GPT-3.5 summaries as input features. The three-stage routing model achieves 0.72 on Task 1 and 0.43 on Task 2, competitive with the fine-tuned baseline despite requiring no task-specific training. Error analysis identifies three confusion clusters: Dodging/Deflection/General, Explicit/Implicit, and a low performance on Partial/half-answer. Contrary to expectations, synthetic data augmentation degraded Task 2 performance, which we attribute to a distributional mismatch between generated and real presidential speech.

1. Introduction

Politicians not always offering the full answers is a well-documented phenomenon in all kinds of political discourse. Empirical studies have shown that politicians provide clear responses to only 39–46% of questions in televised interviews, compared to 70–89% for non-politicians (Bull2003). While this behavior is widely recognised, automated detection of question evasion remains underexplored. The ability to systematically identify and categorise evasive responses has clear implications for political transparency, media analysis, and the broader study of strategic communication.

This paper addresses the task of fine-grained evasion classification as defined by the SemEval-2026 CLARITY - Unmasking Political Question Evasions shared task (Thomas et al.2026), which introduces two subtasks. The first task focuses on three class groupings: *Clear Reply*, *Ambivalent*, and *Clear Non-Reply*. The second task deals with a nine-class taxonomy of response types derived from the three original classes. We focus primarily on Task 2, the nine-class classification subtask, which remains largely unsolved with state-of-the-art macro-F1 scores plateauing around 0.68. The core challenge is severe class imbalance, minority classes such as *Partial/half-answer* and *Clarification* contain fewer than 100 training examples, causing models to systematically ignore them. We address this through three contributions: (1) an error analysis identifying confusion clusters among semantically similar evasion labels, (2) LLM-based data augmentation divided into two approaches: (a) targeting underrepresented classes using GPT-4o-mini and then filtering using GPT-5.4-mini, as for b) we replace existing GPT-3.5 summaries with newer and more informed GPT-5.4-nano-generated summaries to be

used as input features, lastly (3) we introduce a confidence-conditional hierarchical routing pipeline for fine-grained classification.

2. Related Work

Computational analysis of political texts has grown considerably over the past decade, spanning tasks such as ideology classification, stance detection, and topic modelling of political speech (Glavaš et al.2019). Here, the problem of question evasion has been introduced as a distinct classification task, with the QEvasion dataset and a two-level taxonomy of response clarity establishing the first baselines for automated evasion detection in presidential interviews (Thomas et al.2024).

2.1. Fine-Grained Evasion Classification

While high-level clarity classification (Task 1) has reached strong performance with macro-F1 scores of 0.89, fine-grained evasion classification (Task 2) remains a significant challenge. Existing baselines utilizing encoder fine-tuning, such as DeBERTa, RoBERTa, and XLNet, have encountered a performance ceiling around 0.50 for Task 2, largely due to the semantic similarities between categories like *Dodging* and *Deflection*. Even state-of-the-art systems built on large language models (LLMs) plateau at 0.68 macro-F1 (Thomas et al., 2026). Our work differs from these baselines by specifically targeting this performance gap through a systematic application of class imbalance-aware data augmentation.

2.2. Class Imbalance in Text Classification

The QEvasion dataset is very sparse when it comes to certain classes, and very dense when it comes to

other classes. For instance, the *Explicit* category contains 1,051 training samples, while minority classes such as *Partial/half-answer* and *Clarification* contain only 79 and 92 samples, respectively. While standard mitigation strategies like focal loss and oversampling have been applied, they have proven insufficient for resolving the semantic overlaps in political evasion (Thomas et al., 2026). Recent efforts by (Syed and Pedersen, 2026) explored LLM-based augmentation using Gemini 3 and Claude Sonnet 4.5 to improve minority-class recall. We extend this direction by utilizing GPT models for generating and quality filtering new samples. Also, GPT-5.4-nano is utilised to generate extra summary features and replace old GPT-3.5-generated ones.

2.3. Prompt-based approach

The original QEvasion paper also evaluates prompt-based approaches for response clarity classification (Thomas et al., 2024). They test zero-shot, few-shot, and chain-of-thought prompting, in addition to instruction-tuned models. Their prompt-based experiments mainly use GPT-3.5 for zero-shot and chain-of-thought inference, while few-shot experiments are also reported for open-source models such as Llama and Falcon. The prompt templates include the interview segment, the target question, the taxonomy, and an instruction to return the corresponding taxonomy code (Thomas et al., 2024). Our approach follows the same general idea of using an LLM as a prompt-based classifier rather than fine-tuning the model. However, instead of using a single zero-shot, few-shot, or chain-of-thought prompt, we implement a three-stage routing strategy.

3. Experimental Setup

3.1. Dataset

In this study, we utilize the QEvasion dataset (Thomas et al., 2024). Each sample includes the original press conference context (interview_question and interview_answer), a specific target sub-question extracted from the prompt (question), and a set of model-generated textual summaries (gpt3.5_summary). The primary target labels are organized in a hierarchical framework consisting of two classification steps:

- 1. Evasion-Based Clarity (9-class / Low-level):** A granular classification tracking specific, highly nuanced evasive strategies: *Explicit*, *Implicit*, *Dodging*, *Deflection*, *Partial/half-answer*, *General*, *Declining to answer*, *Claims ignorance*, and *Clarification*.
- 2. Direct Clarity (3-class / High-level):** A condensed mapping that hierarchically aggregates the 9 granular strategies into three broader conceptual categories: *Clear Reply*, *Ambivalent*, and *Clear Non-Reply*.




Figure 1: The taxonomy from which the clarity and evasion labels are derived (Thomas et al., 2024)

3.2. Feature setup formats:

We investigate how different contexts affect the model’s ability to spot evasions. The pipeline supports three distinct text concatenation configurations:

- 1. Basic setup:** Merges features of the original dataset: interview_question, interview_answer and question. These are then passed to the model.
- 2. GPT-3.5 summary setup:** Replaces the messy, multi-part interview question context with a structured, LLM-generated summary breakdown of the question.
- 3. GPT-5.4-nano summary setup:** Instead of the GPT-3.5 summary from the dataset, it uses a newly generated summary using the GPT-5.4-nano.

3.3. Model Baseline Setup

For all fine-tuning experiments, we use RoBERTa-base (Liu et al., 2019) as our encoder backbone. Following Thomas et al. (2024), we train exclusively on non-truncated inputs, retaining only samples with fewer than 512 tokens after concatenating the input fields. The reason for this is RoBERTa-base having the upper bound of only 512 input tokens set. This results in approximately 1,700 training samples for the evasion-based clarity and the direct clarity tasks. The only hyperparameter where we decided to stray away from Thomas et al. (2024) was the number of epochs; we decided on fine-tuning for a maximum of fifty epochs instead of five because experimentally we saw a consistent increase in macro F1 scores when letting the models train for more than the originally established five epochs. To mitigate the large number of epochs, we used early stopping with a patience of 3 epochs and a minimum improvement threshold of 0.001 macro F1. Using the method, every model instance converged within 25 epochs. Everything else is set the same as it was in the Thomas et al. (2024) paper; we use the AdamW optimizer with a constant learning rate of 1×10^{-5} . Batch size set to 16 for training and 32 for evaluation. Lastly, we report results on a stratified 20% held-out validation split of the training data, as the official test labels are not publicly released.

3.4. Data Augmentation Setup

To address the severe class imbalance in Task 2, we generate synthetic training examples for the four most underrepresented evasion classes: *Partial/half-answer* (79 samples), *Clarification* (92), *Claims ignorance* (119), and *Declining to answer* (145). Our target is approximately 400 examples per class. It was decided to bring all the classes to a number of samples above the next smallest class size, which was *Deflection* and its 381 samples, but also to mitigate the extra noise that too many generated

examples would bring in. This was also done in hopes of releasing some of the confusion clusters identified with our baselines. These examples are generated using GPT-4o-mini with a structured prompt that includes the label definition derived from the original taxonomy (Thomas et al., 2024), four real demonstrations sampled from the training set, and explicit instructions to produce realistic U.S. presidential interview questions and answers. Generating a total of 1,163 new examples. To improve data quality, generated samples are subsequently filtered using GPT-5.4-mini, which evaluates each example against five criteria: label adherence, realism, evasion-subtlety, absence of cross-label ambiguity, and coherence between the focused sub-question and the full interview question. Samples failing any criterion are discarded. Synthetic examples are added exclusively to the training split. The validation set contains only real examples to ensure unbiased evaluation.

We also tried adding new LLM-generated summaries to improve the NLP classifier’s ability to detect evasion by providing richer, more targeted context for each question-answer pair. Unlike the old summaries, which offered a general overview of the entire exchange, the new summaries were restructured to isolate each sub-question individually and explicitly track how or whether the answer addressed it. To generate these summaries, a reasoning-capable model (GPT-5.4-nano) was used, with the expectation that its stronger inferential capabilities would produce more precise and contextually aware descriptions of answer coverage. The assumption being that more focused, more detailed data would give the classifier stronger signals for identifying the nature of the answer.

3.5. Three-Stage Prompt-Based Routing Model Setup

As an alternative to fine-tuning, we implement a three-stage prompt-based routing model designed to decompose the nine-class evasion classification problem into a series of smaller, more tractable decision steps. The motivation is that broader clarity level categories are more semantically separable than fine-grained evasion labels, making them easier for a language model to distinguish in a zero-shot or few-shot setting.

All stages use a consistent few-shot prompting structure containing: a description of the classification task, target class definitions, decision rules for label assignment, the expected output format, and input field descriptions. Each prompt includes two examples per target class where possible. The input to each stage consists of the original interview question, the specific sub-question, the politician’s answer, a GPT-3.5-generated summary, and a boolean flag indicating whether the question was phrased affirmatively. The model is instructed to return a structured label alongside a confidence score. All stages use GPT-5.4 as the backbone model.

In Stage 1, the model classifies each answer into one of three clarity-level categories: *Clear Reply*, *Ambivalent*, or *Clear Non-Reply*. Examples predicted as *Clear Reply* are assigned the *Explicit* evasion label directly and require no further routing, since this clarity category maps unambigu-

ously to a single evasion type.

In Stage 2, examples classified as *Ambivalent* or *Clear Non-Reply* are routed to specialized prompts that restrict classification to the relevant subset of evasion labels. Each specialized prompt includes a *None of the above* fallback option for cases where the model cannot confidently assign a label within the constrained label set.

In Stage 3, examples assigned *None of the above* in Stage 2 are passed to a final fallback prompt that considers all nine evasion labels simultaneously. This stage handles uncertain or ambiguous cases that the more restricted Stage 2 prompts could not resolve.

3.6. Metrics

The primary metric for this study is the Macro *F1*. Given the class imbalances inherent to political evasion categories, Macro *F1* provides an unweighted arithmetic mean across all classes, ensuring robust performance evaluation independent of class frequencies. To provide an additional breakdown of model behavior across the two tasks we also add Macro Precision, Macro Recall, and Macro Accuracy.

4. Results

Task 1. (3 label-classification) As visible from Table 1, across all models, the best-performing system was the Baseline GPT-3.5 summary variant, achieving a Macro *F1* of 0.7349, followed closely by the Three-stage prompt-based routing model at 0.720. The GPT-5.4 summaries model ranked third with a Macro *F1* of 0.6960, while the single-stage prompt-based routing model reached 0.655. The weakest results were observed for the data augmentation approaches. Added samples (0.6277) and Added samples filtered (0.6276), both falling below the Baseline basic model (0.6435), suggesting that the additional samples did not provide a consistent benefit and may have introduced noise. Overall, the gap between the top model and the Baseline basic is substantial (9 Macro *F1* points), highlighting the importance of richer input representations such as GPT-generated summaries or multi-stage routing.

Task 2. (9 label-classification) As visible from Table 2, performance across all models dropped considerably compared to Task 1, which is expected given the increased label space. The Baseline GPT-3.5 summary model again achieved the highest Macro *F1*, at 0.5253, indicating that GPT-generated summaries provide the most informative features for classification. The Baseline basic followed at 0.4688, while the Three-stage prompt-based routing model reached 0.434 and the GPT-5.4 summaries model scored 0.4561. As in Task 1, the augmentation-based models performed worst, with Added samples at 0.3865 and Added samples filtered at 0.3829, again falling below the Baseline basic. The consistent underperformance of the augmented variants across both tasks suggests that the added training examples, whether unfiltered or filtered, do not align well with the target distribution and may be counterproductive.

Table 1: Model comparison for Task 1 (3 labels).

Model	Macro F1	Accuracy	Recall	Precision
Baseline GPT-3.5 summary	0.7349	0.7450	0.7494	0.7249
Baseline basic	0.6435	0.6575	0.6480	0.6397
Added samples	0.6933	0.7030	0.7087	0.6812
Added samples filtered	0.6833	0.6629	0.6788	0.6891
GPT-5.4-nano summaries	0.6960	0.6992	0.6883	0.7115
Prompt-based routing model	0.6550	0.7119	0.6827	0.6442
Three-stage prompt-based routing	0.7200	0.7500	0.7178	0.7254

Table 2: Model comparison for Task 2 (9 labels).

Model	Macro F1	Accuracy	Recall	Precision
Baseline GPT-3.5 summary	0.5253	0.5826	0.5256	0.5303
Baseline basic	0.4688	0.5547	0.4744	0.4642
Added samples	0.4143	0.4540	0.4139	0.4374
Added samples filtered	0.4017	0.4760	0.3966	0.4472
GPT-5.4-nano summaries	0.4561	0.5055	0.4571	0.4672
Three-stage prompt-based routing	0.4340	0.4958	0.4657	0.4424


5. Analysis

5.1. Summary vs. Basic Input


The most consistent finding across both tasks is the substantial performance gap between the GPT-3.5 summary and basic input setups. On Task 1, the summary-based baseline achieves 0.7349 macro F1 compared to 0.6435 for basic. on Task 2, 0.5253 versus 0.4688. We attribute this to the noisy nature of raw transcripts, which sometimes contain multiple questions merged into one statement and transcription artefacts. GPT-3.5 summaries provide a clean, structured decomposition of the question that the model can reliably use as a classification signal, whereas the basic setup forces the model to extract this signal itself under the 512-token limit, constraining input context.




(a) Summary — Task 1



(b) Summary — Task 2



(c) Basic — Task 1



(d) Basic — Task 2

Figure 2: Confusion matrices for baseline models. Left two: GPT-3.5 summary input. Right two: basic input.

Confusion matrices derived from the baseline outputs

pictured in Figure 2 identify three confusion clusters: *Dodging/Deflection/General*, *Explicit/Implicit*, and a low performance on *Partial/half-answer*.

5.2. Data Augmentation

Contrary to our hypothesis, augmentation consistently degraded performance. Added samples scored 0.6277 on Task 1 and 0.3865 on Task 2, both below the Baseline basic, and quality filtering via GPT-5.4-mini did not improve performance (0.6276 and 0.3829), as it only seemingly limited the amount of training data. We identify a few causes: *data mismatch*, as GPT-4o-mini generates syntactically clean, generic text that does not reflect a person’s “in the moment” answers; and *label boundary degradation*, as the targeted minority classes (*Partial/half-answer* and *Clarification*) have inherently subtle boundaries with adjacent categories that proved difficult to reliably reproduce through generation, as these synthetic examples were the most heavily filtered out after the GPT-5.4-mini filtration step. These findings suggest that LLM-based augmentation for fine-grained evasion detection requires domain-adapted generation models or stronger generation constraints.

5.3. Improved summaries

While the new summary format was designed to offer a more rigorous and structured analysis of question-answer pairs, the NLP model’s stronger performance on the old summary format reveals an important tension between analytical depth and machine-readability. The new summary, though more thorough in decomposing the interview question into discrete sub-parts and tracking answer coverage at a granular level, may have inadvertently introduced noise that the model was not trained to handle. Its verbose, multi-layered structure with labels like “Answer coverage,” “Additional context,” and numbered question parts, departs significantly from the simpler, declarative style the model learned from during training. As a result, the model likely

struggled to extract the clean, bounded signals it relies on to detect evasion patterns.

The old summary, by contrast, presents information in a compact and consistently structured format that closely mirrors the kind of input the model was optimized for. Its numbered, self-contained points map predictably onto features the classifier can evaluate, making it easier for the model to identify whether a question was addressed or deflected. The simplicity that made the old format seem less sophisticated is precisely what made it more compatible with the model’s learned representations.

5.4. Three-stage prompt-based routing model

The results show that the three-stage prompt-based routing model was more effective for broad clarity-level classification than for fine-grained evasion-label classification. The clarity macro F1-score improved from 0.655 with the initial clarity prompt to 0.720 with the full routing model, suggesting that decomposing the task into smaller stages helped the model handle uncertain cases more reliably.

This improvement is expected because the three clarity categories are broader and more semantically distinct than the nine evasion labels. In contrast, evasion-label classification remained more difficult, achieving a macro F1-score of 0.434. This lower performance is likely caused by overlap between similar evasion strategies and by the uneven distribution of labels in the evaluation set, where less frequent labels had less stable per-class scores.

Overall, hierarchical prompting improved clarity-level classification, but fine-grained evasion detection remained challenging due to overlapping label definitions and greater classification complexity.

6. Conclusion

This paper presented our system for the SemEval-2026 CLARITY shared task on political question evasion detection. We explored three complementary directions: error analysis of confusion patterns, LLM-based data augmentation for minority classes, and a three-stage confidence-conditional routing pipeline.

Our results confirm that input representation quality is the dominant factor in model performance. The GPT-3.5 summary setup consistently outperformed all alternatives, achieving macro F1 scores of 0.73 on Task 1 and 0.53 on Task 2, while the newer GPT-5.4-nano summaries, despite being more analytically thorough, underperformed due to their verbose, multi-layered structure introducing noise that proved difficult for the classifier to process.

Data augmentation, contrary to our expectations, consistently degraded performance across both tasks. Synthetic samples generated by GPT-4o-mini lacked the rhetorical sophistication of real presidential evasion, producing examples that were too straightforward to capture the subtle boundary distinctions the model needs to learn.

The three-stage routing pipeline used few-shot prompts to first classify each answer into one of three clarity levels before routing examples to specialized evasion-label prompts, improving clarity-level macro F1 from 0.655 to 0.720. However, fine-grained evasion classification remained harder, reaching only 0.434, indicating that over-

lapping label definitions and class sparsity require more robust modeling strategies beyond hierarchical prompting alone.

Taken together, our findings highlight the difficulty of fine-grained political evasion detection and the gap between intuitive improvements and measurable gains. Future work should focus on domain-adapted generation for augmentation, more classifier-compatible summary formats, and hybrid systems that better handle the semantic overlap inherent to nuanced evasion taxonomies.

References

- Peter Bull. 2003. *The Microanalysis of Political Communication: Claptrap and Ambiguity*. Routledge.
- Goran Glavaš, Federico Nanni, and Simone Paolo Ponzetto. 2019. Computational analysis of political texts: Bridging research efforts across communities. In Preslav Nakov and Alexis Palmer, editors, *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Tutorial Abstracts*, pages 18–23, Florence, Italy, July. Association for Computational Linguistics.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. RoBERTa: A robustly optimized BERT pretraining approach.
- Shujauddin Syed and Ted Pedersen. 2026. Duluth at SemEval-2026 task 6: DeBERTa with LLM-augmented data for unmasking political question evasions. *arXiv preprint arXiv:2604.20168*.
- K. Thomas, G. Filandrianos, M. Lymperaïou, C. Zerva, and G. Stamou. 2024. “I never said that”: A dataset, taxonomy and baselines on response clarity classification. In *Findings of the Association for Computational Linguistics: EMNLP 2024*. Association for Computational Linguistics.
- G. Thomas, C. Tzouvaras, K. Skianis, and A. Voulodimos. 2026. CLARITY: SemEval-2026 task 6 on question evasion and political NLP. In *Proceedings of the 20th International Workshop on Semantic Evaluation (SemEval-2026)*, Barcelona, Spain. Association for Computational Linguistics.